



Sheet 3, starting from Nov 28th, 2021, due Dec 13th, 2021

Introduction and General Comments: We will provide you with

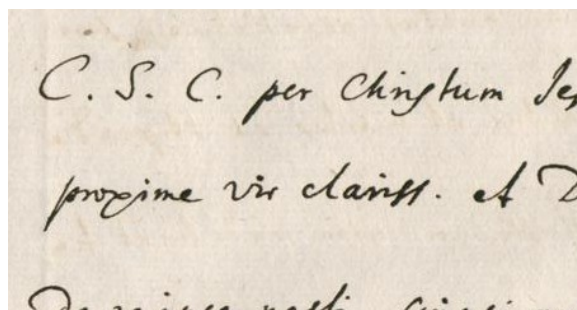
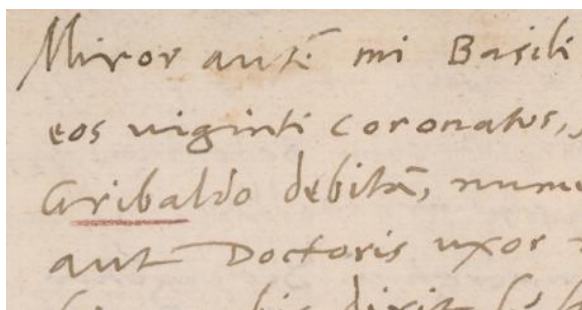
- (1) label files for the training and test set (`icdar17_labels_train.txt` and `icdar17_labels_test.txt`) which contain for each image a class index (i.e., the ground truth writer label), and (2) for each image a `.pkl.gz` file which contains already the local descriptors of each file. Please download the files from (8.7GB):

https://faubox.rrze.uni-erlangen.de/dl/fiY3jdCidNW1AGRHDujPfrBP/icdar17_local_features.zip

Note: in case of quota issues: use your personal ciptmp-folder, i.e. `/proj/ciptmp/<yourlogin>`, it offers you 10GB (no backup). If the folder for your login doesn't exist, create it on your own. (If you cannot see the ciptmp folder, you have to enter it first — it is a shared network drive).

- We also provide in StudOn a code skeleton ('`skeleton.py`'), which we recommend to use.
- Test your algorithms with different parameters and try to get a deeper understanding regarding their behaviors. Show your implementation and your results to one of the advisors.

Exercise 3.1: Writer Identification



节选

Abbildung 1: Image excerpts of the ICDAR2017 Historical WI Dataset: 21-IMG_MAX_966414 (left), 21-IMG_MAX_1006529 (right).

In this exercise, we will implement a small writer identification system. It is going to be evaluated using the ICDAR17 Historical WI benchmark dataset. The dataset is divided into a training set (a.k.a. experiment set) and a test set (a.k.a. benchmark set). Please download the dataset and the code skeleton.

词袋模型

检索

The *bag of (visual) words* model for image recognition/retrieval is going to be used, where each image is represented by a global vector representation, which is formed by encoding the local feature descriptors. The global descriptor is eventually used for the classification/retrieval.

vector of locally aggregated descriptors 局部聚集描述符向量

最大池化

We will use VLAD encoding, which later gets extended by means of *generalized max pooling*. Given T local descriptors $\mathcal{X} = \{\mathbf{x}_t \in \mathbb{R}^D | t = 1, \dots, T\}$ and a dictionary (or codebook) $\mathcal{D} = \{\boldsymbol{\mu}_k \in \mathbb{R}^D | k = 1, \dots, K\}$, the VLAD encoding (i.e. VLAD embedding + sum pooling, see lecture) for one cluster is: $\boldsymbol{\xi}_k(\mathcal{X}) = \sum_{t=1}^T \alpha_k(\mathbf{x}_t)(\mathbf{x}_t - \boldsymbol{\mu}_k)$. The overall encoding is: $\boldsymbol{\xi}(\mathcal{X}) = (\boldsymbol{\xi}_1^T, \dots, \boldsymbol{\xi}_K^T)^T \in \mathbb{R}^{DK}$.

池化pooling的作用：(1) 保留主要特征的同时减少参数和计算量，防止过拟合。(2) invariance(不变性)，这种不变性包括 translation(平移)，rotation(旋转)，scale(尺度)。

dictionary: 字典是另一种可变容器模型，且可存储任意类型对象。包括键key和值value

a) Codebook generation

编码簿

For **computing** the **residuals** $(\mathbf{x}_t - \boldsymbol{\mu}_k)$, we need a codebook to **generate the dictionary** \mathcal{D} . To do so, a representative number of descriptors needs to be selected. Therefore, use a random selection of about 500 000 descriptors from **the training features**.

Compute a codebook (a.k.a. vocabulary/background model) with the help of k -means ($k = 100$). For this operation you can use sklearn's `MiniBatchKMeans`.

b) VLAD **encoding**

With the help of the computed dictionary, we can now **compute the residuals** (see above). They are multiplied with the **indicator/association function** $\alpha_k(\mathbf{x}_i)$, **which returns 1 for the nearest cluster center and 0 otherwise**. Therefore, you can use `cv2.BFMatcher` and `knnMatch`.

Suggestion: compute first the **full association matrix** (i.e. a vector for each descriptor. Each element of the vector represents the association to a cluster. **That means, set all elements of the vector to 0 except the nearest cluster center** ^{assignment}). This can then be used to obtain for each cluster only those descriptors that have this cluster assigned. Then, **compute the residuals with these descriptors** and aggregate them to one global descriptor.

Compute the pairwise distances between all encodings and evaluate the performance in terms of mean average precision (mAP).

c) VLAD **normalization**

In order to improve the representational power of VLAD encoding, we should account for visual burstiness, i.e. **frequent local descriptors that dominate the overall similarity**. This can be achieved by using **power normalization**: $\xi'_i = \text{sign}(\xi_i) \sqrt{|\xi_i|}$, $\forall \{\xi_i\} \in \boldsymbol{\xi}$, where `sign` is +1 if $\xi_i \geq 0$ and -1 otherwise. Compute the pairwise distances between all **normalized encodings** and evaluate the performance in terms of mean average precision (mAP) and compare with the result before.

d) Exemplar classification

The last step of the pipeline is the proposed exemplar classification. Therefore, you have to compute **an individual SVM** for each global **representation (encoding)** of the *test* set using this encoding as **positive** and **all the encodings of the training set as negatives**. (sklearn's `LinearSVC`), a good C value is 1000 (actually this value should be cross-validated with a different validation set), also set `class_weight` to `balanced`. Afterwards, ℓ^2 **normalize the weight vector** (`coef_` parameter of the SVM) and use it as your new global descriptor.

Compute the **pairwise distances between all newly created descriptors** and evaluate the performance in terms of mean average precision (mAP) and compare with the result before.

e) **Bonus: Generalized max pooling**

To further improve the result (unfortunately, this has only very little effect on these descriptors), we now **exchange the sum-pooling part of VLAD encoding with generalized max pooling**. Therefore, you can use sklearn's ridge regression (`Ridge`) with the conjugate gradient solver (`sparse_cg`, note: `fit_intercept` can be set to `False`, `max_iter` to 500, you can try out different regularization values, standard is $\gamma = 1$). You can apply the ridge regression component-wise for each cluster to save some computation time (since the residuals are independent from each other). **The encoding lies then in the `coef_` parameter of the ridge regression classifier.**

Compute the effect of GMP with and without the exemplar classifiers.

Submission Please submit your code `exercise3.py` and a `exercise3.txt` file in which you report the result for each of the experiments.